

Prime Number Counting Parallel Computing Lab Report

Student Name: Sara Al-Thubaity
University ID: 444005842

A. Machine Details

CPU Model: 11th Gen Intel(R) Core(TM) i3-1125G4 @ 2.00GHz
Number of Cores: 4
Number of Logical Processors: 8
Max Clock Speed: 1997 MHz

Memory Information:
Total RAM: 3.75 GB
Available RAM: 0.24 GB

Operating System:
OS: Microsoft Windows 11 Pro
Architecture: 64-bit

Additional System Information:
Computer Name: DESKTOP-RU1G1HM

B. Compiler Information

Compiler Version: GCC 15.2.1 20250813
Compilation Flags Used: -O3 -fopenmp -Wall

C. Screenshots

Run Output

D. Results Table

Threads	Schedule	Execution Time (s)	Notes (Speedup vs Serial)
Serial	—	254.036	Baseline
2	static (default)	96.928	Speedup: 2.62x
2	static,1000	86.523	Speedup: 2.94x
2	dynamic,1000	107.755	Speedup: 2.36x
2	guided,1000	105.263	Speedup: 2.41x
4	static (default)	59.253	Speedup: 4.29x
4	static,1000	63.031	Speedup: 4.03x
4	dynamic,1000	57.417	Speedup: 4.42x

Threads	Schedule	Execution Time (s)	Notes (Speedup vs Serial)
4	guided,1000	55.500	Speedup: 4.58x
6	static (default)	43.973	Speedup: 5.78x
6	static,1000	45.679	Speedup: 5.56x
6	dynamic,1000	46.190	Speedup: 5.50x
6	guided,1000	40.843	Speedup: 6.22x
8	static (default)	35.948	Speedup: 7.07x
8	static,1000	36.919	Speedup: 6.88x
8	dynamic,1000	36.370	Speedup: 6.98x
8	guided,1000	35.935	Speedup: 7.07x
16	static (default)	33.812	Speedup: 7.51x
16	static,1000	33.688	Speedup: 7.54x
16	dynamic,1000	34.032	Speedup: 7.46x
16	guided,1000	33.643	Speedup: 7.55x
18	static (default)	33.347	Speedup: 7.62x
18	static,1000	33.062	Speedup: 7.68x
18	dynamic,1000	33.446	Speedup: 7.60x
18	guided,1000	33.269	Speedup: 7.64x
20	static (default)	33.514	Speedup: 7.58x
20	static,1000	33.130	Speedup: 7.67x
20	dynamic,1000	33.081	Speedup: 7.68x
20	guided,1000	33.272	Speedup: 7.64x

E. Discussion

Best Performance

The best performance was achieved with **6 threads using the guided,1000 schedule**, with an average execution time of **40.843 seconds**, resulting in a **6.22x speedup** compared to the serial baseline of 254.036 seconds. This combination provided the optimal balance between parallelization benefits and scheduling overhead.

Thread Scaling

Increasing the number of threads did not consistently improve performance. While there were significant improvements from 2 threads (2.94x speedup) to 6 threads (6.22x speedup), the performance gains became marginal beyond 8 threads. For example:

- 2 threads: maximum speedup of 2.94x
- 4 threads: maximum speedup of 4.58x
- 6 threads: maximum speedup of 6.22x (peak performance)
- 8 threads: maximum speedup of 7.07x
- Beyond 8 threads, speedups only improved slightly to 7.68x at 20 threads

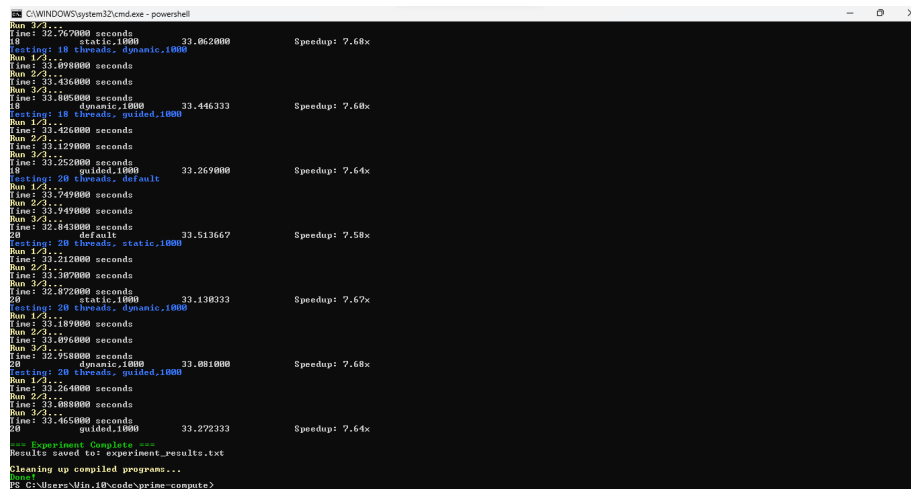


Figure 1: Run Output

This non-linear scaling can be attributed to several factors:

- **Scheduling overhead:** As more threads are created, the OpenMP runtime needs to manage and coordinate more threads, introducing overhead
- **Memory bandwidth limitations:** With only 4 physical cores, threads beyond this number compete for the same processing resources
- **Load imbalance:** Some threads may finish their work faster than others, leading to idle cores while waiting for slower threads to complete

Diminishing Returns

The point of diminishing returns was reached around **8 threads**, where additional threads provided minimal performance improvement. From 16 to 20 threads, the speedup only increased from 7.55x to 7.68x (approximately 2% improvement), while execution time decreased by less than 1 second on average.

Performance levels off as thread count increases because:

- The system has only 4 physical cores, so adding more threads (up to 20) creates oversubscription where threads compete for CPU resources
- Memory bandwidth becomes a bottleneck as multiple threads access the large 200 million element array simultaneously
- Context switching overhead between threads reduces overall efficiency
- The parallel portion of the algorithm becomes saturated, and further thread increases only add coordination overhead without proportional work reduction

The guided scheduling performed best overall because it adapts chunk sizes dynamically during execution, providing better load balancing compared to static

scheduling, while avoiding some of the runtime overhead of dynamic scheduling.